```python
# Import libraries
import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

# Create a simple dataset
data = {
    'Hours_Studied': [2, 4, 5, 7, 8, 10, 12, 13, 15, 16],
    'Attendance': [60, 65, 70, 75, 80, 85, 88, 90, 92, 95],
    'Assignments': [3, 4, 5, 6, 7, 8, 8, 9, 10, 10],
    'Marks': [45, 50, 55, 62, 68, 75, 80, 84, 90, 94]
}

df = pd.DataFrame(data)

# Display dataset
print(df)

# Features and target
X = df[['Hours_Studied', 'Attendance', 'Assignments']]
y = df['Marks']

# Split data into training and testing
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Create and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluation
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

# Display results
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R-squared:", r2)
```

```
   Hours_Studied  Attendance  Assignments  Marks
0              2          60            3     45
1              4          65            4     50
2              5          70            5     55
3              7          75            6     62
4              8          80            7     68
5             10          85            8     75
6             12          88            8     80
7             13          90            9     84
8             15          92           10     90
9             16          95           10     94
MAE: 0.9929078014184375
MSE: 1.0716764750263983
RMSE: 1.0352180809019895
R-squared: 0.9960632292754132
```